

DT²: Decision-Targeted Digital Twins

Anonymous Authors¹

Abstract

A digital twin (DT) is a virtual model of a real-world system that can assist decision-making by simulating scenarios induced by different policies. However, the typical design process of machine learning-based DTs does not optimise for this objective. We prove that, when model capacity is limited, typical DT training paradigms, which minimise one-step transition errors, can produce suboptimal models for ranking sets of policies. We further show that this holds empirically, even with expressive model classes. To address this, we introduce DT², a decision-targeted DT training paradigm. DT² uses off-policy evaluation methods to estimate values of candidate policies on offline data, and encourages the DT to generate rollouts that preserve pairwise policy rankings derived from these proxy ground-truths with an architecture-agnostic loss function. We empirically demonstrate the efficacy of our method across a range of settings and architectures, showing that DT² consistently improves policy ranking and reduces decision regret relative to conventional DT training, both for policies used during training and for unseen policies, while maintaining a good level of raw simulation fidelity.

1. Introduction

A digital twin (DT) is a virtual model of a real-world system that can simulate its dynamics. DTs are appealing in many domains—such as finance (Slepneva et al., 2021), climate science (Voosen, 2020), manufacturing (Rosen et al., 2015), energy (Ismail et al., 2024), agriculture (Purcell & Neubauer, 2023), robotics (Mazumder et al., 2023), and medicine (Katsoulakis et al., 2024)—due to their ability to guide decision-making processes. This manifests with a user deliberating over DT-generated simulations under potential

policies (Tao et al., 2018; Corral-Acero et al., 2020), to determine the best policy or establish a preference ordering. Much of the DT literature highlights this critical use case, impressing that DTs must “*aid decision making*” (Wagg et al., 2020), and “[*inform*] decisions that realize value” (National Academy of Engineering and National Academies of Sciences, Engineering, and Medicine, 2024). To enable such human-in-the-loop decision support, DTs must generate realistic simulations that allow for user interrogation and verification, while ensuring the learned dynamics lead to correct rankings over relevant policies.

In machine learning (ML)-based DTs, however, this goal of decision support is often overlooked. Typically, ML-based DTs involve a model that is trained to minimise a measure of one-step transition error, such as mean squared error (MSE) or negative log likelihood (NLL), across all variables and time points on a dataset of observed state-action trajectories (San et al., 2023; Kuang et al., 2024; Holt et al., 2024; Makarov et al., 2024; Canzaniello et al., 2024; Amad et al., 2025). While this can lead to good simulation fidelity, such direction-, variable-, and timepoint-agnostic objectives do not necessarily correlate with the ability to correctly rank policies. Optimising with respect to them can therefore misalign DTs with their primary downstream purpose, misusing model capacity and under-emphasising variables or temporal slices that are especially crucial for distinguishing between high- and low-value policies.

For example, consider a medical DT, used to help clinicians treat septic shock. Despite large data availabilities in ICU settings, clinicians typically focus on only a few key indicators to make timely decisions (Pickering et al., 2013), like blood pressure and lactate levels for sepsis. A typical DT, trained with a variable-agnostic simulation loss, considers non-critical dimensions as equivalent to such key indicator variables. By wasting model capacity on largely irrelevant features, DTs can provide suboptimal decision support, jeopardising patient health. For a visual example of how simulation training can lead to poor decisions due to being direction-agnostic, see Figure 1.

A core re-framing is necessary to emphasise that, for human-in-the-loop decision support, a DT need not be a perfect replica of previously observed reality. As such, DT construction should move away from such context-agnostic

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.




Figure 1. We visualise two diabetes DTs, $\hat{\Phi}_1$ (top) and $\hat{\Phi}_2$ (bottom), under treatment plans with **high**, **low**, and **adaptive** insulin levels. $\hat{\Phi}_1$ has lower MSE from the true rollouts, yet, considering the ‘value’ of a diabetes policy as the time it spends in the optimal glucose region, it leads to an incorrect preference ordering (**high insulin** $\approx$ **low insulin** $>$ **adaptive**). In contrast, $\hat{\Phi}_2$ leads to a correct preference ordering (**adaptive** $>$ **high insulin** $\approx$ **low insulin**), despite its higher MSE, as its simulations traverse the correct reward regions.

optimisation, towards what is best for decision support.

To resolve this, we introduce DT²: a **decision-targeted DT** training paradigm that balances between seeking fidelity in general, and in those parts of the transition distribution that are most relevant for ranking candidate policies. Since, in general, ground-truth rankings amongst candidate policies are not known *a priori*, we estimate policy values using off-policy evaluation (OPE) methods. OPE methods can estimate scalar policy values using offline data (Fu et al., 2021) but, distinct from DTs, they offer little in the way of transparency or interpretability, limiting their application in high stakes domains (Rudin, 2019). As such, we use OPE estimates to construct proxy ground-truth pairwise rankings amongst candidate policies, and introduce a differentiable ranking loss function, that operates alongside standard simulation fidelity objectives, to encourage DTs to generate simulations that preserve these proxy rankings. By doing so, we distil the value-awareness of OPE into the interpretable, generative structure of DTs. A tunable hyperparameter, λ , modulates the trade-off between raw simulation accuracy and reconstruction of OPE policy rankings. In summary, our contributions are:

1. We prove that standard DT training is theoretically suboptimal for decision-making purposes when the DT hypothesis space is limited, establishing a necessity to introduce a decision-aware training mechanism (§3).

2. We propose DT², an architecture-agnostic decision-targeted DT training paradigm that aligns generation with proxy policy rankings (§4). We use a smooth approximation of Kendall’s rank correlation coefficient (Kendall, 1938) (§4.1) to enable supervision with pairwise comparisons of OPE estimates (§4.2).
3. We empirically demonstrate the alignment gap predicted by our theory (§6.1), and show that this persists even with highly expressive hypothesis spaces (§6.2). Across a suite of continuous control tasks, DT² reduces decision regret by a factor of nearly four, at an MSE cost of only 19%, demonstrating better downstream decision-making capacity. Furthermore, we show that this advantage generalises to out-of-distribution settings, achieving better preference orderings amongst policies never seen during training (§6.3).

2. Formalism

We now formalise the problem of decision support that we aim to address with DTs.

2.1. The Decision-Making Goal

Let a real-world system be modelled as a Markov Decision Process (MDP) denoted by $\mathcal{M} := (\mathcal{X}, \mathcal{A}, \Phi, r, \gamma)$. Here, $\mathcal{X} \subseteq \mathbb{R}^{d_{\mathcal{X}}}$ is a $d_{\mathcal{X}}$ -dimensional state space, $\mathcal{A} \subseteq \mathbb{R}^{d_{\mathcal{A}}}$ is a $d_{\mathcal{A}}$ -dimensional action space, $\Phi : \mathcal{X} \times \mathcal{A} \rightarrow \mathcal{P}(\mathcal{X})$ is the transition distribution, where $\mathcal{P}(\mathcal{X})$ denotes the set of probability measures over $\mathcal{X}$, $r : \mathcal{X} \times \mathcal{A} \times \mathcal{X} \rightarrow \mathbb{R}$ is a bounded reward function, and $\gamma \in [0, 1)$ is a discount factor. At time t , given the current state $\mathbf{x}_t \in \mathcal{X}$ and action $\mathbf{a}_t \in \mathcal{A}$, the next state is sampled according to $\mathbf{x}_{t+1} \sim \Phi(\cdot | \mathbf{x}_t, \mathbf{a}_t)$.

A policy $\pi : \mathcal{X} \rightarrow \mathcal{P}(\mathcal{A})$ can interact with $\mathcal{M}$. Applying π from an initial state $\mathbf{x}_0$ yields a trajectory $\tau = (\mathbf{x}_0, \mathbf{a}_0, \mathbf{x}_1, \mathbf{a}_1, \dots) \sim (\Phi, \pi)$. We define the value of a policy π in $\mathcal{M}$ as the expected sum of discounted rewards:

$$J(\pi) := \mathbb{E}_{\tau \sim (\Phi, \pi)} \left[\sum_{t=0}^{\infty} \gamma^t r(\mathbf{x}_t, \mathbf{a}_t, \mathbf{x}_{t+1}) \right]. \quad (1)$$

In decision-making processes, a user may want to interrogate trajectories induced by each π in some set Π , as well as establish a preference ordering over Π . This is defined by the relation $\succ$, where for any pair $\pi_i, \pi_j \in \Pi$, $\pi_i \succ \pi_j \iff J(\pi_i) > J(\pi_j)$. Alternatively, a reduced goal may be to identify the optimal policy π^* :

$$\pi^* := \arg \max_{\pi \in \Pi} J(\pi). \quad (2)$$

Generating such trajectories, and determining such policy rankings exactly is generally impossible, as it would require repeated experimentation on the real-world system.

2.2. Digital Twins as Practical Surrogates

A DT can serve as a practical surrogate for the real system, involving a parametrised model $\hat{\Phi}_\theta$ that is trained to approximate Φ . Given a policy π , the DT can simulate a trajectory from $\mathbf{x}_0$, denoted $\hat{\tau} = (\hat{\mathbf{x}}_0, \mathbf{a}_0, \hat{\mathbf{x}}_1, \mathbf{a}_1, \dots) \sim (\hat{\Phi}_\theta, \pi)$. We denote the DT-estimated value of π as:

$$\hat{J}(\pi; \theta) := \mathbb{E}_{\hat{\tau} \sim (\hat{\Phi}_\theta, \pi)} \left[\sum_{t=0}^{\infty} \gamma^t r(\hat{\mathbf{x}}_t, \mathbf{a}_t, \hat{\mathbf{x}}_{t+1}) \right], \quad (3)$$

which leads to a DT preference ordering defined by $\succsim$ such that $\pi_i \succsim \pi_j \iff \hat{J}(\pi_i; \theta) > \hat{J}(\pi_j; \theta)$, and a DT-optimal policy $\hat{\pi}^*$:

$$\hat{\pi}^* := \arg \max_{\pi \in \Pi} \hat{J}(\pi; \theta). \quad (4)$$

3. Misalignment of Conventional Digital Twin Training

Conventional ML-based DT training paradigms involve a training dataset $\mathcal{D} = \{(\mathbf{x}, \mathbf{a}, \mathbf{x}')_i\}_{i=1}^n$ of n observed transitions collected from the real-world system, or related systems. The optimal model parameters θ^* are typically found by minimising a measure of simulation error, $\mathcal{L}_{\text{sim}}(\theta)$, over $\mathcal{D}$. There are a few possible instantiations of $\mathcal{L}_{\text{sim}}(\theta)$, with the most notable being NLL for probabilistic models:

$$\mathcal{L}_{\text{sim}}^{\text{NLL}}(\theta) := \mathbb{E}_{(\mathbf{x}, \mathbf{a}, \mathbf{x}') \sim \mathcal{D}} \left[-\log \hat{\Phi}_\theta(\mathbf{x}' | \mathbf{x}, \mathbf{a}) \right], \quad (5)$$

or MSE when focusing on the expected trajectory or assuming constant variance:

$$\mathcal{L}_{\text{sim}}^{\text{MSE}}(\theta) := \mathbb{E}_{(\mathbf{x}, \mathbf{a}, \mathbf{x}') \sim \mathcal{D}} \left[\|\mathbf{x}' - \hat{\mu}_\theta(\mathbf{x}, \mathbf{a})\|^2 \right], \quad (6)$$

where $\hat{\mu}_\theta(\mathbf{x}, \mathbf{a}) = \mathbb{E}_{\mathbf{x}' \sim \hat{\Phi}_\theta(\cdot | \mathbf{x}, \mathbf{a})}[\mathbf{x}']$. Such loss functions are a poor proxy for the ultimate goal of establishing preference orderings over policies.

Theorem 3.1. *Let Φ be the transition distribution of a real-world system and $\mathcal{F} = \{\hat{\Phi}_\theta | \theta \in \Theta\}$ be a DT hypothesis class such that $\Phi \notin \mathcal{F}$ and*

$$\min_{\theta \in \Theta} \mathbb{E}_{(x, a) \sim \mathcal{D}} \left[D_{\text{KL}}(\Phi(\cdot | x, a) \| \hat{\Phi}_\theta(\cdot | x, a)) \right] > 0. \quad (7)$$

Let $\hat{\Phi}_{\theta^}$ be the $\mathcal{L}_{\text{sim}}$ -optimal DT, i.e.*

$$\theta^* \in \arg \min_{\theta \in \Theta} \mathcal{L}_{\text{sim}}(\theta). \quad (8)$$

Then, there exists a reward function r and a set of policies Π such that $\hat{\Phi}_{\theta^}$ induces a suboptimal preference ordering on Π according to r .*

Proof. See Appendix A. $\square$




Figure 2. When $\Phi \notin \mathcal{F}$, the $\mathcal{L}_{\text{sim}}$ -optimal DT, $\hat{\Phi}_{\theta^*}$, may land in a region of low decision quality (grey areas). With decision-targeted DT training, we wish to find $\hat{\Phi}_{\theta'}$ that resides in a better decision region (green areas), at a small simulation fidelity cost.

From Theorem 3.1 we can see that when a DT cannot exactly model its target system, optimising $\mathcal{L}_{\text{sim}}$ will not generally prioritise decision-making. In equation-based DTs, or small-model DTs, Theorem 3.1 can directly apply. Given that collecting large data for certain real-world systems can be difficult, and that DTs may need to constantly update alongside their counterpart system (Amad et al., 2025), such small models can be common. Furthermore, even when expressive models are used, the complex dynamics or inherent partial observability of real-world systems may still result in non-covering hypothesis spaces. Moreover, under a further assumption on $\mathcal{F}$, we can see that better possible models exist for decision-making.

Definition 3.2. Let $(x_0, a_0) \in \text{supp}(\mathcal{D})$ and $G \subseteq \mathcal{X}$ be such that $\Phi(G | x_0, a_0) \neq \hat{\Phi}_{\theta^*}(G | x_0, a_0)$. Define

$$p := \Phi(G | x_0, a_0), \quad \hat{p} := \hat{\Phi}_{\theta^*}(G | x_0, a_0). \quad (9)$$

Assume, w.l.o.g., that $p > \hat{p}$. Then, $\mathcal{F}$ allows *local improvement at θ^** if there exists $\theta' \in \Theta$ such that

$$\hat{p}' := \hat{\Phi}_{\theta'}(G | x_0, a_0) > \hat{p}. \quad (10)$$

This definition applies to $\mathcal{F}$ when it contains an alternative model to the $\mathcal{L}_{\text{sim}}$ -optimal model whose predictions on a local subset of transitions are closer to the true transition distribution.

Theorem 3.3. *Assume the conditions of Theorem 3.1 and that $\mathcal{F}$ allows local improvement at θ^* . Then there exists a reward function r , a set of policies Π , and $\theta' \in \Theta$ such that $\hat{\Phi}_{\theta'}$ induces strictly better preference ordering on Π according to r than $\hat{\Phi}_{\theta^*}$.*

Proof. See Appendix A. $\square$

Theorem 3.3 outlines our theoretical motivation to search for a better training paradigm for DTs than optimising $\mathcal{L}_{\text{sim}}$, that will better converge to the parameters θ' that permit good preference ordering (Figure 2). We elaborate on the

scenarios when these theorems directly apply, and empirically validate them, in §6.1. Moreover, we hypothesise that, even when $\Phi \in \mathcal{F}$, optimising $\mathcal{L}_{\text{sim}}$ will poorly converge to an optimal decision-making model, and we show empirical evidence for this in §6.2 and §6.3.

4. DT²: Decision-Targeted Digital Twin Training

While §3 motivates our pursuit to train decision-targeted DTs, there exist some practical hurdles to this. First, comparing a ground-truth preference ordering $\succ$ to a DT-estimated $\hat{\succ}$ with standard ranking metrics involves indicator functions, i.e. checking $\mathbb{I}(\pi_i \succ \pi_j) \iff \mathbb{I}(\pi_i \hat{\succ} \pi_j) \forall \pi_i, \pi_j \in \Pi$, which are not differentiable. We address this by using a differentiable ranking loss to instil decision-targeting (§4.1). Second, the ground truth $\succ$ is almost never available, and indeed this is in part the reason for constructing a DT in the first place. So we derive proxy preference orderings to target, using OPE-based value estimates (§4.2). Finally, to avoid the need to generate and backpropagate through long rollouts when optimising $\hat{\succ}$ during training, we employ a value-bootstrapped scheme (§4.3). This forms DT², our method for DT training that encourages DTs to allocate capacity toward modelling dynamics that are critical for decision support.

4.1. Differentiable Ranking Loss

To instil decision-targeting, we define a ranking objective based on a smooth approximation of Kendall’s rank correlation coefficient (Kendall, 1938). Let $\Delta_{ij} = J(\pi_i) - J(\pi_j)$ be the ground-truth value difference between two policies, and $\hat{\Delta}_{ij}(\theta) = \hat{J}(\pi_i; \theta) - \hat{J}(\pi_j; \theta)$ be the corresponding DT-estimated value difference. Kendall’s correlation measures the concordance of signs between these differences:

$$\frac{1}{|\mathcal{C}|} \sum_{(i,j) \in \mathcal{C}} \text{sign}(\Delta_{ij}) \cdot \text{sign}(\hat{\Delta}_{ij}(\theta)), \quad (11)$$

where $\mathcal{C}$ is the set of policy pairs in a candidate set Π . To make this suitable for gradient-based optimisation, we replace the sign functions with hyperbolic tangents, and minimise the complement as our ranking loss:

$$\mathcal{L}_{\text{rank}}(\theta) := 1 - \frac{1}{|\mathcal{C}|} \sum_{(i,j) \in \mathcal{C}} \tanh\left(\frac{\Delta_{ij}}{\alpha}\right) \cdot \tanh\left(\frac{\hat{\Delta}_{ij}(\theta)}{\alpha}\right) \quad (12)$$

where the temperature α governs the trade-off between faithfulness to the discrete Kendall correlation and gradient smoothness. Smaller α produces sharper, more sign-like curves and larger α yields more stable, smooth gradients during optimisation.

A key benefit of this formulation is that $\tanh(\Delta_{ij}/\alpha)$ acts

as an implicit weighting mechanism. When the reference evaluator is indifferent between two policies ($\Delta_{ij} \approx 0$), the term $\tanh(\Delta_{ij}/\alpha)$ approaches zero, causing the gradient contribution for that pair to vanish. Consequently, the model is encouraged to ignore noise in policy separation, allocating capacity toward distinguishing between high-regret pairs. We empirically compare this formulation to alternative hinge-based and listwise losses in Appendix D.

4.2. Off-Policy Evaluation for Proxy Preferences

Since ground-truth policy values $J(\pi)$ are inaccessible in offline settings, we derive proxy Δ_{ij} targets using principles from OPE (§5.2). Specifically, we utilize Fitted Q-Evaluation (FQE) (Le et al., 2019) to estimate the action-value function $Q^\pi(\mathbf{x}, \mathbf{a})$ for each $\pi \in \Pi$. $Q^\pi(\mathbf{x}, \mathbf{a})$ estimates the expected long-term return of executing action $\mathbf{a}$ in state $\mathbf{x}$ and following π thereafter:

$$Q^\pi(\mathbf{x}, \mathbf{a}) := \mathbb{E}_{\tau \sim (\Phi, \pi)} \left[\sum_{t=0}^{\infty} \gamma^t r_t \mid \mathbf{x}_0 = \mathbf{x}, \mathbf{a}_0 = \mathbf{a} \right] \quad (13)$$

FQE approximates this with a parameterised function Q_ψ^π trained on $\mathcal{D}$ by iteratively minimising the expected Bellman error:

$$\mathcal{L}_{Q^\pi}(\psi) := \mathbb{E}_{\mathcal{D}} \left[\left(Q_\psi^\pi(\mathbf{x}, \mathbf{a}) - \left(r + \gamma Q_\psi^\pi(\mathbf{x}', \pi(\mathbf{x}')) \right) \right)^2 \right] \quad (14)$$

where Q_ψ^π is a target network. The parameters $\bar{\psi}$ are frozen copies of ψ that are updated periodically to stabilise the bootstrapping target. Once converged, the scalar value of π can be estimated by averaging Q_ψ^π over the initial states,

$$J_{\text{FQE}}(\pi) \approx \mathbb{E}_{\mathbf{x}_0 \sim \mathcal{D}, \mathbf{a} \sim \pi(\mathbf{x}_0)} [Q_\psi^\pi(\mathbf{x}_0, \mathbf{a})]. \quad (15)$$

These estimates allow us to construct the proxy pairwise rankings targets $\Delta_{ij}^{\text{FQE}} = J_{\text{FQE}}(\pi_i) - J_{\text{FQE}}(\pi_j)$ that are used in place of Δ_{ij} in Eq. 12.

We select FQE amongst the plethora of OPE methods given its simplicity and strong performance on OPE benchmarks (Fu et al., 2021; Voloshin et al., 2021). Furthermore, while FQE is prone to systematic bias—particularly for out-of-distribution policies—due to consistent overestimation of Q-values, these errors will tend to cancel out during pairwise comparisons. This makes FQE-based ranking more robust than its scalar value estimates, as relative orderings are preserved even when absolute value magnitudes are biased.

4.3. Bootstrapping Digital Twin Value Estimates

To compute the DT-estimated ranking targets $\hat{\Delta}_{ij}(\theta)$, the model must simulate trajectories $\hat{\tau} \sim (\hat{\Phi}_\theta, \pi)$ (Eq. 3). However, backpropagating through long rollouts to optimize $\mathcal{L}_{\text{rank}}$ is computationally expensive and prone to vanishing

or exploding gradients (Bengio et al., 1994; Pascanu et al., 2013). To mitigate this, we employ a bootstrapping scheme that truncates DT rollouts at a fixed horizon H and approximates the remaining return using a frozen Q_{ψ}^{π} . The bootstrapped DT-value estimate for policy π is defined as:

$$\hat{J}_H(\pi; \theta) := \mathbb{E}_{\hat{\tau} \sim (\hat{\Phi}_{\theta}, \pi)} \left[\sum_{t=0}^{H-1} \gamma^t r_t + \gamma^H Q_{\psi}^{\pi}(\hat{\mathbf{x}}_H, \mathbf{a}_H) \right]. \quad (16)$$

By using $\hat{J}_H$ to compute $\hat{\Delta}_{ij}(\theta)$, we reduce the cost of calculating $\mathcal{L}_{\text{rank}}$, and improve the its gradient stability.

To ensure DT² still permits generation of faithful trajectories and enables human-in-the-loop decision support, the ultimate training objective combines a $\mathcal{L}_{\text{rank}}$ with $\mathcal{L}_{\text{sim}}$:

$$\mathcal{L}_{\text{DT}^2}(\theta) := (1 - \lambda)\mathcal{L}_{\text{sim}}(\theta) + \lambda\mathcal{L}_{\text{rank}}(\theta). \quad (17)$$

The hyperparameter $\lambda \in [0, 1]$ modulates the trade-off between optimising for raw simulation fidelity and decision-targeting. $\lambda \rightarrow 0$ recovers standard DT training, while $\lambda \rightarrow 1$ yields a pure decision-targeted model. We investigate its effects in §6.4. By shaping the loss landscape with Δ_{ij}^{FOE} , DT² distills the value-awareness and decision-making power of ‘black-box’ OPE methods into the interpretable, generative structure of a DT.

5. Related Works

5.1. Digital Twins

DTs have been researched since their use in 1960s aerospace engineering (Allen, 2021), where well-known physical laws could fully describe system dynamics. While manually-defined DTs remain in use (Laubenbacher et al., 2024), their creation requires significant domain expertise. Methods like genetic programming (Koza, 1994), SINDy (Brunton et al., 2016), and PySR (Cranmer, 2023) aim to automate the discovery of such mechanistic models, but they are often outperformed by more expressive deep learning approaches.

Deep learning DTs can approximate system dynamics from longitudinal data with a variety of model architectures, including standard MLP approaches (Holt et al., 2024), as well as tailored sequential models like RNNs (Elman, 1990; Graves, 2012), transformers (Vaswani et al., 2017), and pre-trained (Amad et al., 2025) or fine-tuned LLMs (Makarov et al., 2024). Neural ODEs (Chen et al., 2018; Dupont et al., 2019; Alvarez et al., 2020) further extend deep learning capabilities to modelling continuous-time dynamics.

Hybrid approaches combine mechanistic and deep learning methods to improve sample efficiency and generalisation. These range from models that combine neural components with well-defined physical equations (Raissi et al., 2019; Kuang et al., 2024) to those that integrate learnable com-

ponents into known model structures (Qian et al., 2021; Takeishi & Kalousis, 2021).

5.2. Off-Policy Evaluation

OPE methods are used in offline reinforcement learning (RL) settings to estimate $J(\pi_{\text{target}})$ from an offline dataset collected from some behavioural policy $\pi_{\text{behavioral}}$. Popular approaches include direct methods (Lagoudakis & Parr, 2003; Ernst et al., 2005; Munos & Szepesvári, 2008; Le et al., 2019), which learn and use $Q^{\pi_{\text{target}}}$ to estimate $J(\pi_{\text{target}})$, importance sampling methods (Precup et al., 2000; Hallak & Mannor, 2017; Liu et al., 2018; Xie et al., 2019; Nachum et al., 2019; Uehara et al., 2020; Zhang et al., 2020a;b), which use importance weights to correct for the difference in densities between $\pi_{\text{behavioral}}$ and π_{target} to estimate $J(\pi_{\text{target}})$, and doubly robust methods (Jiang & Li, 2016; Thomas & Brunskill, 2016), which combine direct and importance sampling methods.

While OPE methods can be directly used for decision support, a primary limitation is their lack of interpretability, as they function as ‘black-boxes’ that produce scalar value estimates that are difficult to verify. Because of this, in our work we use OPE methods as a supervisory signal to distil decision-awareness into a more interpretable final model—a DT that generates trajectories a user can inspect, understand, and reject if they seem implausible. Moreover, DTs provide a natural avenue for encoding inductive biases about system dynamics, such as physics-based mechanistic equations, which can increase generalisation in regions not well-covered by offline data (Holt et al., 2024).

5.3. Model-Based Reinforcement Learning

A DT is similar to the dynamics model in model-based RL (MBRL) (Sutton, 1991), however MBRL uses this model to learn new policies, rather than our objective to rank a fixed set of policies. Some MBRL research has a similar motivation to ours, recognising that global accuracy is unnecessary (Lambert et al., 2020; Grimm et al., 2020). Proposed solutions include policy-aware methods (Eysenbach et al., 2022; Wang et al., 2023; Ma et al., 2023) that favour areas of the state-action space that are traversed by the current policy, and value-aware methods (Farahmand et al., 2017; Farahmand, 2018; Grimm et al., 2020; Farquhar et al., 2021; Voelcker et al., 2022), which are more similar to DT², as they focus model capacity in areas important to determine the current policy value. However, unlike DT², these methods do not operate with a fixed, known set of policies, which may all have distinct value functions and visitation distributions. Furthermore, since the models in MBRL are purely a means to an end, they do not focus on maintaining simulation fidelity, for interpretability purposes, which is essential for human-in-the-loop decision support.

6. Empirical Investigation

We now empirically validate the efficacy of DT^2 . We cover the limited capacity setting that relates to Theorems 3.1 and 3.3 in §6.1, before showing that their takeaways hold in the expressive setting as well in §6.2 and §6.3. Finally, we investigate the effects of λ in $\mathcal{L}_{\text{DT}^2}$ in §6.4. We provide detailed experimental set-ups in Appendix B.

6.1. Restricted Hypothesis Spaces

To empirically demonstrate Theorems 3.1 and 3.3, we construct three toy scenarios which satisfy the required assumptions, where $\Phi \notin \mathcal{F}$ and $\mathcal{F}$ allows local improvement at θ^* . N.B. in the following, we assume access to ground-truth policy values during DT^2 training, for simplicity. In the more realistic settings, from §6.2 onward, we demonstrate the full DT^2 pipeline, with FQE-derived proxy rankings.

6.1.1. LIMITED NEURAL CAPACITY

We construct a two-dimensional system $x_t = [x_{d,t}, x_{c,t}]^\top$ that decomposes into independent ‘decoy’ and ‘critical’ components, with an associated scalar action a_t . x_t evolves according to $x_{d,t+1} = M \sin(x_{d,t})$ and $x_{c,t+1} = x_{c,t} + \epsilon a_t$, and we set M and ϵ such that $\text{Var}(x_{d,t}) \gg \text{Var}(x_{c,t})$. We set $\mathcal{F}$ as the family of three layer MLPs with a single-neuron bottleneck as the middle layer. Forcing an internal one-dimensional representation prevents any model from simultaneously representing the independent evolution of $x_{d,t}$ and $x_{c,t}$, guaranteeing $\Phi \notin \mathcal{F}$.

We consider Π with two constant-action policies, $\pi_{\text{high}} = 1.0$ and $\pi_{\text{low}} = 0.9$, and a reward $r(\mathbf{x}_t, a_t, \mathbf{x}_{t+1}) = x_{c,t}$ such that $\pi_{\text{high}} \succ \pi_{\text{low}}$. Across 20 seeds, we generate $\mathcal{D}$ by drawing $a_t \sim U[-1, 1]$, and train DTs with $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ and $\mathcal{L}_{\text{DT}^2}$. Because $x_{d,t}$ dominates the variance in $\mathcal{D}$, the $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ -trained DT allocates its significant capacity, and consequently achieves only a 20% success rate across seeds in identifying π_{high} as the optimal policy via model-estimated policy values. In contrast, $\mathcal{L}_{\text{DT}^2}$ encourages the bottlenecked model to identify that x_c drives value differences, despite its small magnitude, resulting in 100% ranking accuracy.

6.1.2. MISSPECIFIED MECHANISTIC MODEL

We now consider a mechanistic setting, where the assumed functional form of the dynamics is incorrect. The true one-dimensional system has cubic action-dependent dynamics, $x_{t+1} = x_t + (a_t - a_t^3) + \xi$, where $\xi \sim \mathcal{N}(0, 0.01)$, and we set $\mathcal{F}$ as the family of linear models, $x_{t+1} = x_t + \beta a_t$, such that $\Phi \notin \mathcal{F}$. We collect $\mathcal{D}$ by drawing $a_t \sim U[-1.5, 1.5]$.

We consider Π containing three constant-action policies $\pi_1 = -0.58$, $\pi_2 = 0$, $\pi_3 = +0.58$, and a reward $r(x_t, a_t, x_{t+1}) = x_t$ such that $\pi_3 \succ \pi_2 \succ \pi_1$. We show the mechanistic models learnt by optimising $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ and $\mathcal{L}_{\text{DT}^2}$,

and the true dynamics in Figure 3. For the action range in $\mathcal{D}$, the cubic term $-a_t^3$ tends to dominate, creating a broadly negative correlation between action and state change $\delta = x_{t+1} - x_t$, and the $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ -trained DT learns this negative slope, consequently ordering $\pi_1 \succ \pi_2 \succ \pi_3$. DT^2 , on the other hand, recognises that, in the local decision-critical region for Π , the linear term dominates, sacrificing global fit to learn a positive slope, recovering the true preference ordering.




Figure 3. Mechanistic models learnt via $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ and $\mathcal{L}_{\text{DT}^2}$, compared to the true Φ . Black dots represent policies in Π .

6.1.3. PARTIALLY OBSERVED DYNAMICS

Finally, we investigate a system with partially observable dynamics, where an unobserved latent variable z_t dictates the transition distribution, ensuring $\Phi \notin \mathcal{F}$ even with an expressive hypothesis space of three-layer MLPs. We consider a one-dimensional system governed by $x_{t+1} = x_t + a_t + z_t$, where $z_t \sim \mathcal{N}(0, 25)$ dominates, and set Π with two constant-action policies, $\pi_{\text{high}} = 0.1$ and $\pi_{\text{low}} = -0.1$, and a reward $r(x_t, a_t, x_{t+1}) = x_t$ such that $\pi_{\text{high}} \succ \pi_{\text{low}}$. Across 20 seeds, we collect a small $\mathcal{D}$ by drawing $a_t \sim U[-1, 1]$.

Minimising $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ encourages learning the conditional expectation $\mathbb{E}[x_{t+1} | x_t, a_t]$ by averaging over the unobserved latent z_t . While a_t and z_t are independent in the generative process, the high variance of z_t can cause, with finite data, spurious correlations to emerge where actions coincide with opposing latent values. Consequently, the empirical conditional mean often exhibits a slope flip, causing the $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ model to learn a negative relationship between a_t and x_{t+1} to fit the noise in $\mathcal{D}$. As such, it elicits the correct preference ordering on Π only 65% of the time. In contrast, DT^2 utilises the ranking loss to distinguish π_{high} from π_{low} , better conserving the true positive slope of the action, achieving 100% ranking accuracy.

6.2. Expressive Hypothesis Spaces

We now wish to test our hypothesis that the takeaways from Theorems 3.1 and 3.3 persist even when $\mathcal{F}$ is expressive and could contain Φ . We utilise standard Mujoco (Todorov et al., 2012) and Gymnasium (Towers et al., 2024) continuous control environments (Pendulum, Lunarlander, Hopper,




Figure 4. Regret and Spearman’s correlation for preference orderings from $\mathcal{L}_{\text{sim}}^{\text{NLL}}$ (hashed) and $\mathcal{L}_{\text{DT}^2}$ (solid) DTs across base architectures in six continuous control environments. Adapted VaGraM method shown in grey. We report averages over 10 seeds, with 95% CIs.

Walker, Cheetah, Ant) to act as the ‘real-world systems’ Φ . These environments span a range of complexities, from 4- to 35-dimensional state-action spaces. For each Φ , we define a candidate set Π of six policies taken from evenly-spaced checkpoints of a PPO (Schulman et al., 2017) training run. We then generate $\mathcal{D}$ by deploying each $\pi \in \Pi$ in Φ for 10,000 steps, and learn each Q_{ψ}^{π} using FQE on $\mathcal{D}$.

We consider three DT variants. 1) an $\mathcal{L}_{\text{sim}}^{\text{NLL}}$ -trained DT represents the prevailing context-agnostic approach. 2) a value-aware MBRL-inspired approach, adapted to our setting. VaGraM (Voelcker et al., 2022) is a method designed for active policy learning, which weights transition errors by the gradient of the value function being learned concurrently. To adapt this to our setting, with multiple, frozen candidate policies Π , we train the DT using an MSE loss weighted by the average squared gradient norm of all Q_{ψ}^{π} -functions:

$$\mathbb{E}_{\mathcal{D}} \left[\|\mathbf{x}' - \mu_{\theta}(\mathbf{x}, \mathbf{a})\|^2 \cdot \frac{1}{|\Pi|} \sum_{\pi \in \Pi} \|\nabla_{\mathbf{x}'} Q_{\psi}^{\pi}(\mathbf{x}', \pi(\mathbf{x}'))\|^2 \right]. \quad (18)$$

This weights the importance of transitions by how sensitive policy returns are to errors in those regions. 3) we train a model via DT^2 , with $\lambda = 0.1$ and $\alpha = 1$.

We use a number of expressive base architectures for the $\mathcal{L}_{\text{sim}}^{\text{NLL}}$ and $\mathcal{L}_{\text{DT}^2}$ DTs, to show the generality of our method. These include using a Transformer (Vaswani et al., 2017), Gated Recurrent Unit (GRU) (Cho et al., 2014), multi-layer perceptron (MLP), residual network (ResNet) (He et al., 2016), and neural ODE (Chen et al., 2018) to parametrise the mean and variance of a multivariate Gaussian distribution. For the adapted VaGraM DT, we use an MLP architecture.

To evaluate the DT-derived preference orderings, $\succ$, against the ground-truth $\succ$, we report the Spearman’s rank correlation (Spearman, 1904), to measure global ranking alignment, and the decision regret, to measure the value gap between π^* and $\hat{\pi}^*$, in Figure 4. $\mathcal{L}_{\text{DT}^2}$ DTs outperform their $\mathcal{L}_{\text{sim}}^{\text{NLL}}$ counterpart in 90% of the architecture-environment-metric couplings, and there is no single environment or architecture that DT^2 does not outperform in, on average. Furthermore, the best DT^2 model outperforms VaGraM in each environment. The average regret, across all environments and architectures, is 6.21 using $\mathcal{L}_{\text{sim}}^{\text{NLL}}$, and 1.61 using $\mathcal{L}_{\text{DT}^2}$, demonstrating a near $4\times$ improvement. Similarly, for Spearman’s correlation, $\mathcal{L}_{\text{sim}}^{\text{NLL}}$ DTs average 0.47 while $\mathcal{L}_{\text{DT}^2}$ DTs average 0.67, demonstrating a 44% improvement.

In terms of simulation fidelity, we do see that there is a necessary trade-off to achieve this increased preference ordering performance. The average transition MSE on a test set using $\mathcal{L}_{\text{DT}^2}$ is 1.32, while it is 1.11 using $\mathcal{L}_{\text{sim}}^{\text{NLL}}$, demonstrating a 19% performance decrease for DT^2 in this regard. These results validate our core thesis: models optimised solely for trajectory reconstruction often misallocate capacity to dynamics irrelevant for policy differentiation, even when using expressive hypothesis classes. Through explicit decision-targeted training, DT^2 recovers significantly better preference orderings, at a modest MSE cost.

6.3. Case Study – Cancer Treatment Digital Twin

We now investigate a medical case study, where Π is comprised of human-defined treatment plans. We use a Cancer environment from DTR-Bench (Luo et al., 2024) that simulates the progression of cancer with metastasis under chemotherapy and radiotherapy, based on the mathematical

model of (Ghaffari et al., 2016), with a nine-dimensional state-action space and a reward function that balances tumor reduction against treatment toxicity. We define five distinct potential treatment plans: 1) no treatment, 2) fractionated radiotherapy, 3) metronomic chemotherapy, 4) adaptive combined therapy, and 5) aggressive combined therapy. We pursue a similar training pipeline as §6.2, using the ResNet architecture as the backbone for the $\mathcal{L}_{\text{sim}}^{\text{NLL}}$ and $\mathcal{L}_{\text{DT}^2}$ DTs, since it tended to perform best in the previous environments.

We report the results in Table 1. Once again, using $\mathcal{L}_{\text{DT}^2}$ leads to a significantly better regret and Spearman’s rank correlation than $\mathcal{L}_{\text{sim}}^{\text{NLL}}$, indicating substantially better potential for clinical decision support. This is heightened by the fact that the MSE cost is only 2% here, minimally impacting human interrogation of DT-generated simulations.

Table 1. Regret and Spearman’s correlation for preference orderings from $\mathcal{L}_{\text{sim}}^{\text{NLL}}$ and $\mathcal{L}_{\text{DT}^2}$ DTs in Cancer environment. Averaged over 5 seeds with standard errors.

Loss	Spearman ($\uparrow$)	Regret ($\downarrow$)	MSE ($\downarrow$)
$\mathcal{L}_{\text{sim}}^{\text{NLL}}$	0.10 (0.08)	59.56 (0.83)	0.329 (0.01)
$\mathcal{L}_{\text{DT}^2}$	0.64 (0.10)	26.96 (11.04)	0.334 (0.01)

Furthermore, we test preference orderings across 11 policies that are unseen during training. These include six PPO checkpoints, obtained as in §6.2, and five further human-defined policies: 1) pulsed chemotherapy, 2) hypofractionated radiotherapy, 3) aggressive combined therapy followed by maintenance chemo, 4) alternating modality treatment, and 5) dose-escalating combined treatment. From Table 2 we see that DT^2 performs best again, indicating that decision-targeting is not overfitting to the policies observed during training, but encourages learning dynamics that assist in general decision-making according to r .

Table 2. Regret and Spearman’s correlation for preference orderings from $\mathcal{L}_{\text{sim}}^{\text{NLL}}$ and $\mathcal{L}_{\text{DT}^2}$ DTs in Cancer environment across unseen policies. Averaged over 10 seeds with standard errors.

Loss	Spearman ($\uparrow$)	Regret ($\downarrow$)
$\mathcal{L}_{\text{sim}}^{\text{NLL}}$	0.40 (0.04)	95.88 (18.23)
$\mathcal{L}_{\text{DT}^2}$	0.50 (0.06)	49.27 (17.63)

6.4. The Effect of λ in $\mathcal{L}_{\text{DT}^2}$

A key choice in $\mathcal{L}_{\text{DT}^2}$ is setting λ , which decides between preserving the FQE-derived policy ranks and the simulation fidelity. We investigate the performance of $\mathcal{L}_{\text{DT}^2}$ DTs across λ values in the six continuous control environments from §6.2, using the same set-ups and, again, using the ResNet architecture for each DT. Figure 5 plots the Spearman’s correlation and test MSE versus λ . N.B. that the $\lambda = 0$ setting is equivalent to training with $\mathcal{L}_{\text{sim}}^{\text{NLL}}$. As expected, we

observe a consistent increase in ranking performance and decrease in simulation fidelity as λ increases. Practitioners can therefore effectively use λ to navigate to their desired point along this trade-off during training. It seems that ranking performance experiences diminishing returns at larger λ values, so we generally suggest setting a relatively small $\lambda > 0$ for most use cases.




Figure 5. Spearman’s correlation and test MSE for $\mathcal{L}_{\text{DT}^2}$ DTs in six continuous control environments across λ levels. Averaged over 5 seeds, with 95% CIs.

7. Discussion

This work challenges the prevailing practice of maximising simulation fidelity in DT construction. We introduced DT^2 , a training paradigm that aligns DTs with their primary downstream purpose: decision guidance. By distilling OPE-based policy rankings into generative DT models, DT^2 induces significantly better preference orderings over candidate policies at a modest simulation fidelity cost. We demonstrate that these benefits are robust across diverse environments, architectures and simulation losses (see $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ results in Appendix C), and that the trade-off between decision-making and fidelity can be navigated via λ .

Limitations. DT^2 depends on the quality of its proxy targets; biased or high-variance FQE estimates resulting from poor data coverage in $\mathcal{D}$ may be replicated by the DT. To address this, future work could explore improving the general efficacy or sample-efficiency of OPE estimators, or investigate some uncertainty-aware training methods, to neglect the signal from $\mathcal{L}_{\text{rank}}$ when the proxy values are likely poor. Additionally, DT^2 introduces computational overhead compared to standard training, requiring pre-trained Q -functions and H -step backpropagation through time. Potential strategies to mitigate this could involve using only a small representative subset of policies in Π during training, or using adaptive bootstrapping horizons.

Impact Statement

This paper presents work whose goal is to advance the field of machine learning. There are many potential societal

consequences of our work, none of which we feel must be specifically highlighted here.

References

- Allen, B. D. Digital twins and living models at NASA. In *Digital Twin Summit*, 2021. URL <https://ntrs.nasa.gov/citations/20210023699>.
- Alvarez, V. M. M., Roşca, R., and Fălcuţescu, C. G. Dynode: Neural ordinary differential equations for dynamics modeling in continuous control. *arXiv preprint arXiv:2009.04278*, 2020. URL <https://arxiv.org/abs/2009.04278>.
- Amad, H., Astorga, N., and van der Schaar, M. Continuously updating digital twins using large language models. In *International Conference on Machine Learning*. PMLR, 2025. URL <https://openreview.net/pdf?id=jk3TjZAHem>.
- Bengio, Y., Simard, P., and Frasconi, P. Learning long-term dependencies with gradient descent is difficult. *IEEE transactions on neural networks*, 5(2):157–166, 1994.
- Brunton, S. L., Proctor, J. L., and Kutz, J. N. Discovering governing equations from data by sparse identification of nonlinear dynamical systems. *Proceedings of the National Academy of Sciences*, 113(15):3932–3937, 2016. doi: 10.1073/pnas.1517384113. URL <https://www.pnas.org/doi/abs/10.1073/pnas.1517384113>.
- Canzaniello, M., Amitrano, S., Prezioso, E., Giampaolo, F., Cuomo, S., and Piccialli, F. Leveraging digital twins and generative ai for effective urban mobility management. In *2024 IEEE Cyber Science and Technology Congress (CyberSciTech)*, pp. 146–153. IEEE, 2024.
- Chen, R. T., Rubanova, Y., Bettencourt, J., and Duvenaud, D. K. Neural ordinary differential equations. *Advances in neural information processing systems*, 31, 2018. URL https://proceedings.neurips.cc/paper_files/paper/2018/file/69386f6bb1dfed68692a24c8686939b9-Paper.pdf.
- Cho, K., Van Merriënboer, B., Gulcehre, C., Bahdanau, D., Bougares, F., Schwenk, H., and Bengio, Y. Learning phrase representations using rnn encoder-decoder for statistical machine translation. *arXiv preprint arXiv:1406.1078*, 2014.
- Corral-Acero, J., Margara, F., Marciniak, M., Roderio, C., Loncaric, F., Feng, Y., Gilbert, A., Fernandes, J. F., Bukhari, H. A., Wajdan, A., Martinez, M. V., Santos, M. S., Shamohammdi, M., Luo, H., Westphal, P., Leeson, P., DiAchille, P., Gurev, V., Mayr, M., Geris, L., Pathmanathan, P., Morrison, T., Cornelussen, R., Prinzen, F., Delhaas, T., Doltra, A., Sitges, M., Vigmond, E. J., Zacur, E., Grau, V., Rodriguez, B., Remme, E. W., Niederer, S., Mortier, P., McLeod, K., Potse, M., Pueyo, E., Bueno-Orovio, A., and Lamata, P. The ‘digital twin’ to enable the vision of precision cardiology. *European Heart Journal*, 41(48):4556–4564, 03 2020. ISSN 0195-668X. doi: 10.1093/eurheartj/ehaa159. URL <https://doi.org/10.1093/eurheartj/ehaa159>.
- Cranmer, M. Interpretable machine learning for science with pysr and symbolicregression.jl. *arXiv preprint arXiv:2305.01582*, 2023.
- Dupont, E., Doucet, A., and Teh, Y. W. Augmented neural odes. *Advances in neural information processing systems*, 32, 2019. URL https://proceedings.neurips.cc/paper_files/paper/2019/file/21be9a4bd4f81549a9d1d241981cec3c-Paper.pdf.
- Elman, J. L. Finding structure in time. *Cognitive Science*, 14(2):179–211, 1990. ISSN 0364-0213. doi: [https://doi.org/10.1016/0364-0213\(90\)90002-E](https://doi.org/10.1016/0364-0213(90)90002-E). URL <https://www.sciencedirect.com/science/article/pii/036402139090002E>.
- Ernst, D., Geurts, P., and Wehenkel, L. Tree-based batch mode reinforcement learning. *Journal of Machine Learning Research*, 6, 2005.
- Eysenbach, B., Khazatsky, A., Levine, S., and Salakhutdinov, R. R. Mismatched no more: Joint model-policy optimization for model-based rl. *Advances in Neural Information Processing Systems*, 35:23230–23243, 2022.
- Farahmand, A.-m. Iterative value-aware model learning. *Advances in Neural Information Processing Systems*, 31, 2018.
- Farahmand, A.-m., Barreto, A., and Nikovski, D. Value-aware loss function for model-based reinforcement learning. In *Artificial Intelligence and Statistics*, pp. 1486–1494. PMLR, 2017.
- Farquhar, G., Baumli, K., Marinho, Z., Filos, A., Hessel, M., van Hasselt, H. P., and Silver, D. Self-consistent models and values. *Advances in Neural Information Processing Systems*, 34:1111–1125, 2021.
- Fu, J., Norouzi, M., Nachum, O., Tucker, G., Wang, Z., Novikov, A., Yang, M., Zhang, M. R., Chen, Y., Kumar, A., et al. Benchmarks for deep off-policy evaluation. *arXiv preprint arXiv:2103.16596*, 2021.

- Ghaffari, A., Bahmaie, B., and Nazari, M. A mixed radio-therapy and chemotherapy model for treatment of cancer with metastasis. *Mathematical methods in the applied sciences*, 39(15):4603–4617, 2016.
- Graves, A. Long short-term memory. *Supervised Sequence Labelling with Recurrent Neural Networks*, pp. 37–45, 2012. URL <https://doi.org/10.1007/978-3-642-24797-2>.
- Grimm, C., Barreto, A., Singh, S., and Silver, D. The value equivalence principle for model-based reinforcement learning. *Advances in neural information processing systems*, 33:5541–5552, 2020.
- Hallak, A. and Mannor, S. Consistent on-line off-policy evaluation. In *International Conference on Machine Learning*, pp. 1372–1383. PMLR, 2017.
- He, K., Zhang, X., Ren, S., and Sun, J. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770–778, 2016.
- Holt, S., Liu, T., and van der Schaar, M. Automatically learning hybrid digital twins of dynamical systems. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024. URL <https://openreview.net/pdf?id=SOsiObSdU2>.
- Ismail, F. B., Al-Faiz, H., Hasini, H., Al-Bazi, A., and Kazem, H. A. A comprehensive review of the dynamic applications of the digital twin technology across diverse energy sectors. *Energy Strategy Reviews*, 52:101334, 2024.
- Jiang, N. and Li, L. Doubly robust off-policy value evaluation for reinforcement learning. In *International conference on machine learning*, pp. 652–661. PMLR, 2016.
- Katsoulakis, E., Wang, Q., Wu, H., Shahriyari, L., Fletcher, R., Liu, J., Achenie, L., Liu, H., Jackson, P., Xiao, Y., et al. Digital twins for health: a scoping review. *NPJ Digital Medicine*, 7(1):77, 2024. URL <https://doi.org/10.1038/s41746-024-01073-0>.
- Kendall, M. G. A new measure of rank correlation. *Biometrika*, 30(1-2):81–93, 1938.
- Koza, J. R. Genetic programming as a means for programming computers by natural selection. *Statistics and computing*, 4:87–112, 1994. URL <https://doi.org/10.1007/BF00175355>.
- Kuang, K., Dean, F., Jedlicki, J. B., Ouyang, D., Philipakis, A., Sontag, D., and Alaa, A. Med-real2sim: Non-invasive medical digital twins using physics-informed self-supervised learning. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024. URL <https://openreview.net/pdf?id=xCUXJqQySD>.
- Lagoudakis, M. G. and Parr, R. Least-squares policy iteration. *Journal of machine learning research*, 4(Dec): 1107–1149, 2003.
- Lambert, N., Amos, B., Yadan, O., and Calandra, R. Objective mismatch in model-based reinforcement learning. *Proceedings of Machine Learning Research* vol, 120:1–15, 2020.
- Laubenbacher, R., Mehrad, B., Shmulevich, I., and Trayanova, N. Digital twins in medicine. *Nature Computational Science*, 4(3):184–191, 2024. URL <https://doi.org/10.1038/s43588-024-00607-6>.
- Le, H., Voloshin, C., and Yue, Y. Batch policy learning under constraints. In *International Conference on Machine Learning*, pp. 3703–3712. PMLR, 2019.
- Liu, Q., Li, L., Tang, Z., and Zhou, D. Breaking the curse of horizon: Infinite-horizon off-policy estimation. *Advances in neural information processing systems*, 31, 2018.
- Luo, Z., Zhu, M., Liu, F., Li, J., Pan, Y., Zhou, J., and Zhu, T. Dtr-bench: An in silico environment and benchmark platform for reinforcement learning based dynamic treatment regime. *CoRR*, 2024.
- Ma, Y. J., Sivakumar, K., Yan, J., Bastani, O., and Jayaraman, D. Learning policy-aware models for model-based reinforcement learning via transition occupancy matching. In *Learning for Dynamics and Control Conference*, pp. 259–271. PMLR, 2023.
- Makarov, N., Bordukova, M., Rodriguez-Esteban, R., Schmich, F., and Menden, M. P. Large language models forecast patient health trajectories enabling digital twins. *medRxiv*, pp. 2024–07, 2024. URL <https://doi.org/10.1101/2024.07.05.24309957>.
- Mazumder, A., Sahed, M. F., Tasneem, Z., Das, P., Badal, F. R., Ali, M. F., Ahamed, M. H., Abhi, S. H., Sarker, S. K., Das, S. K., et al. Towards next generation digital twin in robotics: Trends, scopes, challenges, and future. *Heliyon*, 9(2), 2023.
- Munos, R. and Szepesvári, C. Finite-time bounds for fitted value iteration. *Journal of Machine Learning Research*, 9 (5), 2008.
- Nachum, O., Chow, Y., Dai, B., and Li, L. Dualdice: Behavior-agnostic estimation of discounted stationary distribution corrections. *Advances in neural information processing systems*, 32, 2019.

- National Academy of Engineering and National Academies of Sciences, Engineering, and Medicine. *Foundational Research Gaps and Future Directions for Digital Twins*. The National Academies Press, Washington, DC, 2024. ISBN 978-0-309-70042-9. doi: 10.17226/26894. URL <https://doi.org/10.17226/26894>.
- Pascanu, R., Mikolov, T., and Bengio, Y. On the difficulty of training recurrent neural networks. In *International conference on machine learning*, pp. 1310–1318. Pmlr, 2013.
- Pickering, B. W., Gajic, O., Ahmed, A., Herasevich, V., and Keegan, M. T. Data utilization for medical decision making at the time of patient admission to icu. *Critical care medicine*, 41(6):1502–1510, 2013.
- Precup, D., Sutton, R. S., and Singh, S. P. Eligibility traces for off-policy policy evaluation. In *Proceedings of the Seventeenth International Conference on Machine Learning*, pp. 759–766, 2000.
- Purcell, W. and Neubauer, T. Digital twins in agriculture: A state-of-the-art review. *Smart Agricultural Technology*, 3: 100094, 2023.
- Qian, Z., Zame, W., Fleuren, L., Elbers, P., and van der Schaar, M. Integrating expert odes into neural odes: pharmacology and disease progression. *Advances in Neural Information Processing Systems*, 34:11364–11383, 2021. URL <https://openreview.net/pdf?id=tDqef76wFa0>.
- Raffin, A., Hill, A., Gleave, A., Kanervisto, A., Ernestus, M., and Dormann, N. Stable-baselines3: Reliable reinforcement learning implementations. *Journal of Machine Learning Research*, 22(268):1–8, 2021. URL <http://jmlr.org/papers/v22/20-1364.html>.
- Raissi, M., Perdikaris, P., and Karniadakis, G. E. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving non-linear partial differential equations. *Journal of Computational physics*, 378:686–707, 2019. URL <https://doi.org/10.1016/j.jcp.2018.10.045>.
- Rosen, R., von Wichert, G., Lo, G., and Bettenhausen, K. D. About the importance of autonomy and digital twins for the future of manufacturing. *IFAC-PapersOnLine*, 48(3):567–572, 2015. ISSN 2405-8963. doi: <https://doi.org/10.1016/j.ifacol.2015.06.141>. URL <https://www.sciencedirect.com/science/article/pii/S2405896315003808>. 15th IFAC Symposium on Information Control Problems in Manufacturing.
- Rudin, C. Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead. *Nature machine intelligence*, 1(5):206–215, 2019.
- San, O., Pawar, S., and Rasheed, A. Decentralized digital twins of complex dynamical systems. *Scientific Reports*, 13(1):20087, 2023.
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Slepneva, T., Chernysheva, M., Zaitseva, K., et al. Impact of digital twin technology on the financial performance of corporations. *European Proceedings of Social and Behavioural Sciences*, 2021. URL <http://dx.doi.org/10.15405/epsbs.2021.04.02.145>.
- Spearman, C. The proof and measurement of association between two things. *The American Journal of Psychology*, 15(1):72–101, 1904.
- Sutton, R. S. Dyna, an integrated architecture for learning, planning, and reacting. *ACM Sigart Bulletin*, 2(4):160–163, 1991.
- Takeishi, N. and Kalousis, A. Physics-integrated variational autoencoders for robust and interpretable generative modeling. *Advances in Neural Information Processing Systems*, 34:14809–14821, 2021. URL <https://openreview.net/pdf?id=0p0gt1Pn2Gv>.
- Tao, F., Cheng, J., Qi, Q., Zhang, M., Zhang, H., and Sui, F. Digital twin-driven product design, manufacturing and service with big data. *The International Journal of Advanced Manufacturing Technology*, 94:3563–3576, 2018. URL <https://doi.org/10.1007/s00170-017-0233-1>.
- Thomas, P. and Brunskill, E. Data-efficient off-policy policy evaluation for reinforcement learning. In *International conference on machine learning*, pp. 2139–2148. PMLR, 2016.
- Todorov, E., Erez, T., and Tassa, Y. Mujoco: A physics engine for model-based control. In *2012 IEEE/RSJ international conference on intelligent robots and systems*, pp. 5026–5033. IEEE, 2012.
- Towers, M., Kwiatkowski, A., Terry, J., Balis, J. U., De Cola, G., Deleu, T., Goulão, M., Kallinteris, A., Krimmel, M., KG, A., et al. Gymnasium: A standard interface for reinforcement learning environments. *arXiv preprint arXiv:2407.17032*, 2024.

- Uehara, M., Huang, J., and Jiang, N. Minimax weight and q-function learning for off-policy evaluation. In *International Conference on Machine Learning*, pp. 9659–9668. PMLR, 2020.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., and Polosukhin, I. Attention is all you need. *Advances in neural information processing systems*, 30, 2017.
- Voelcker, C., Liao, V., Garg, A., and Farahmand, A.-m. Value gradient weighted model-based reinforcement learning. *arXiv preprint arXiv:2204.01464*, 2022.
- Voloshin, C., Le, H. M., Jiang, N., and Yue, Y. Empirical study of off-policy policy evaluation for reinforcement learning. In *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track (Round 1)*, 2021.
- Voosen, P. Europe is building a ‘digital twin’ of earth to revolutionize climate forecasts. *Science*, 2020. URL <https://doi.org/10.1126/science.abf0687>.
- Wagg, D., Worden, K., Barthorpe, R., and Gardner, P. Digital twins: state-of-the-art and future directions for modeling and simulation in engineering dynamics applications. *ASCE-ASME Journal of Risk and Uncertainty in Engineering Systems, Part B: Mechanical Engineering*, 6(3): 030901, 2020.
- Wang, X., Wongkamjan, W., Jia, R., and Huang, F. Live in the moment: Learning dynamics model adapted to evolving policy. In *International Conference on Machine Learning*, pp. 36470–36493. PMLR, 2023.
- Xie, T., Ma, Y., and Wang, Y.-X. Towards optimal off-policy evaluation for reinforcement learning with marginalized importance sampling. *Advances in neural information processing systems*, 32, 2019.
- Zhang, R., Dai, B., Li, L., and Schuurmans, D. Gendice: Generalized offline estimation of stationary values. *arXiv preprint arXiv:2002.09072*, 2020a.
- Zhang, S., Liu, B., and Whiteson, S. Gradientdice: Rethinking generalized offline estimation of stationary values. In *International Conference on Machine Learning*, pp. 11194–11203. PMLR, 2020b.

A. Theorem Proofs

Herein we provide the formal proofs for Theorem 3.1 and Theorem 3.3.

A.1. Proof of Theorem 3.1

Proof. By the theorem assumption, $\mathbb{E}_{(x,a) \sim \mathcal{D}}[D_{\text{KL}}(\Phi \| \hat{\Phi}_{\theta^*})] > 0$. This implies the existence of a reachable state-action pair $(x_0, a_0) \in \text{supp}(\mathcal{D})$ and a measurable set $G \subseteq \mathcal{X}$ where the transition probabilities differ:

$$\Phi(G \mid x_0, a_0) \neq \hat{\Phi}_{\theta^*}(G \mid x_0, a_0). \quad (19)$$

Let $p := \Phi(G \mid x_0, a_0)$ and $\hat{p} := \hat{\Phi}_{\theta^*}(G \mid x_0, a_0)$. Without loss of generality, assume $p > \hat{p}$.

Consider a decision scenario with horizon $H = 1$ (or, alternatively, where $r_t = 0$ for $t > 0$). Assume $|\mathcal{A}| \geq 2$, such that there be a second action $a_1 \in \mathcal{A}$ available at x_0 . We define a bounded reward function $r : \mathcal{X} \times \mathcal{A} \times \mathcal{X} \rightarrow [0, 1]$:

$$r(x, a, x') := \mathbb{I}\{x = x_0\} (\mathbb{I}\{a = a_0\} \mathbb{I}\{x' \in G\} + \mathbb{I}\{a = a_1\} c), \quad (20)$$

where $c \in (0, 1)$ is a constant chosen such that $\hat{p} < c < p$.

Let $\Pi = \{\pi_0, \pi_1\}$ be a set of deterministic policies where $\pi_0(x_0) = a_0$ and $\pi_1(x_0) = a_1$. The policy values under the true system Φ are:

$$J(\pi_0) = p, \quad J(\pi_1) = c. \quad (21)$$

Since $p > c$, the true preference ordering is $\pi_0 \succ \pi_1$. However, under $\hat{\Phi}_{\theta^*}$ the estimated policy values are:

$$\hat{J}(\pi_0; \theta^*) = \hat{p}, \quad \hat{J}(\pi_1; \theta^*) = c. \quad (22)$$

Since $\hat{p} < c$, the model induces preference ordering $\pi_1 \hat{\succ} \pi_0$. Thus, $\hat{\Phi}_{\theta^*}$ yields an incorrect preference ordering. $\square$

A.2. Proof of Theorem 3.3

To prove the existence of a better model in $\mathcal{F}$, we rely on the assumption of local improvement: that $\mathcal{F}$ contains an alternative model whose predictions on a local subset of transitions move in the correct direction relative to $\hat{\Phi}_{\theta^*}$.

Proof. From the proof of Theorem 3.1, we have $p > \hat{p}$ and we consider the same one-step decision task with reward

$$r(x, a, x') := \mathbb{I}\{x = x_0\} (\mathbb{I}\{a = a_0\} \mathbb{I}\{x' \in G\} + \mathbb{I}\{a = a_1\} c),$$

and policy set $\Pi = \{\pi_0, \pi_1\}$ where $\pi_0(x_0) = a_0$ and $\pi_1(x_0) = a_1$.

By assuming Definition 3.2, $\hat{p}' > \hat{p}$. Choose any constant

$$c \in (\hat{p}, \min\{p, \hat{p}'\}). \quad (23)$$

This interval is non-empty because $\hat{p}' > \hat{p}$ and $p > \hat{p}$.

True ordering. As before, $J(\pi_0) = p$ and $J(\pi_1) = c$, hence $\pi_0 \succ \pi_1$ since $p > c$ by (23).

Ordering under θ^* . Under $\hat{\Phi}_{\theta^*}$, $\hat{J}(\pi_0; \theta^*) = \hat{p}$ and $\hat{J}(\pi_1; \theta^*) = c$, hence $\pi_1 \hat{\succ}_{\theta^*} \pi_0$ since $c > \hat{p}$.

Ordering under θ' . Under $\hat{\Phi}_{\theta'}$, $\hat{J}(\pi_0; \theta') = \hat{p}'$ and $\hat{J}(\pi_1; \theta') = c$, hence $\pi_0 \hat{\succ}_{\theta'} \pi_1$ since $\hat{p}' > c$ by (23).

Therefore, on the same decision problem (r, Π) , $\hat{\Phi}_{\theta^*}$ induces an incorrect preference ordering, whereas $\hat{\Phi}_{\theta'}$ induces the correct preference ordering. That is, $\hat{\Phi}_{\theta'}$ yields a strictly better preference ordering on Π than $\hat{\Phi}_{\theta^*}$. $\square$

B. Experimental Details

We provide detailed set-ups for all experiments run here. Furthermore, we will release code upon acceptance.

B.1. Limited Neural Capacity Experiment

System Dynamics. The environment consists of a 2-dimensional state space $x_t = [x_{d,t}, x_{c,t}]^\top$ and a 1-dimensional action space $a_t \in \mathbb{R}$. The transitions are defined as:

$$x_{d,t+1} = M \sin(x_{d,t}) \quad (24)$$

$$x_{c,t+1} = x_{c,t} + \epsilon a_t \quad (25)$$

where $M = 500$, $\epsilon = 0.01$. The horizon is $H = 20$. The reward function is $r(x_t, a_t, x_{t+1}) = x_{c,t}$.

Data Generation. We collected a dataset $\mathcal{D}$ of 3,000 episodes using a behavior policy that samples actions uniformly $a \sim \mathcal{U}(-1, 1)$.

Model Architecture. We utilise an MLP with an explicit information bottleneck.

- **Input:** State (2) + Action (1).
- **Encoder:** Linear(3 $\rightarrow$ 32), Tanh, Linear(32 $\rightarrow$ 1).
- **Bottleneck:** The 1-dimensional output of the encoder represents the compressed state representation z .
- **Decoder:** Linear(2 $\rightarrow$ 32) (taking z and a), Tanh, Linear(32 $\rightarrow$ 2).

Training. We trained for 30 epochs with a learning rate of 1×10^{-3} using the Adam optimiser and $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ as the simulation loss. For DT², we used $\lambda = 0.99$ and $\alpha = 1.0$. The ranking targets were computed using the ground truth value difference between π_{high} and π_{low} .

B.2. Misspecified Mechanistic Model Experiment

System Dynamics. The true system is a 1D scalar system governed by cubic control inputs:

$$x_{t+1} = x_t + (a_t - a_t^3) + \xi_t, \quad \xi_t \sim \mathcal{N}(0, 0.1) \quad (26)$$

Hypothesis Space. The model is constrained to be linear in actions:

$$\hat{x}_{t+1} = x_t + \beta a_t \quad (27)$$

Here, β is the only learnable parameter.

Training Data. The dataset consists of 5,000 transitions where actions are sampled uniformly from: $a \sim \mathcal{U}(-1.5, 1.5)$.

Policies. We evaluate three constant policies:

1. $\pi_1 : a = -0.58$
2. $\pi_2 : a = 0.0$
3. $\pi_3 : a = +0.58$

The optimal policy is π_3 .

Training. We trained for 2000 epochs with a learning rate of 0.01 using the Adam optimiser and $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ as the simulation loss. For DT², we used $\lambda = 0.9$ and $\alpha = 1.0$. The ranking targets were computed using the ground truth policy values.

B.3. Partially Observed Dynamics Experiment

System Dynamics. The true system is a 1-dimensional linear system governed by an unobserved latent variable z_t with high variance. The transitions are defined as:

$$x_{t+1} = x_t + a_t + z_t, \quad z_t \sim \mathcal{N}(0, 25) \quad (28)$$

where the state $x_t \sim \mathcal{N}(0, 1)$ and the action $a_t \in [-1, 1]$. The large standard deviation of the latent variable ($\sigma_z = 5$) relative to the action magnitude mimics a scenario where unobserved factors dominate the transition dynamics.

Data Generation. To emphasise data scarcity, we collected a small dataset $\mathcal{D}$ of $N = 100$ transitions. Actions in the dataset were sampled uniformly: $a \sim \mathcal{U}(-1, 1)$.

Model Architecture. We utilise an MLP with a single hidden layer to approximate the dynamics $x_{t+1} \approx \hat{\Phi}_\theta(x_t, a_t)$.

- **Input:** State x_t and Action a_t (Dimension: 2).
- **Hidden:** Linear($2 \rightarrow 16$), ReLU activation.
- **Output:** Linear($16 \rightarrow 1$) representing the predicted next state $\hat{x}_{t+1}$.

Policies. We evaluate the model’s ability to rank two constant-action policies:

1. $\pi_{\text{high}} = 0.01$
2. $\pi_{\text{low}} = -0.01$

Training. We trained for 200 epochs with a learning rate of 1×10^{-2} using the Adam optimiser. We used $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ as the simulation loss. For DT², we used $\lambda = 0.9$ and a temperature $\alpha = 0.1$. The ranking targets were computed using the known ground truth policy values.

B.4. Continuous Control Experiments

B.4.1. ENVIRONMENTS

We use six continuous control environments from the Gymnasium library (Towers et al., 2024) (<https://github.com/Farama-Foundation/Gymnasium>, MIT license), utilising the MuJoCo physics engine (Todorov et al., 2012). These environments were selected to cover a diverse range of complexities, spanning state-action space dimensionalities from 4 to 35. The specific details for each environment are as follows:

- **Pendulum-v1** A classic control problem where the goal is to swing up and balance a pendulum in an inverted position from a random start.
 - $\mathcal{X}$: $d_{\mathcal{X}} = 3$, including the cosine and sine of the pendulum’s angle, and its angular velocity.
 - $\mathcal{A}$: $d_{\mathcal{A}} = 1$, the scalar joint torque applied to the pendulum.
- **LunarLanderContinuous-v3**: A rocket trajectory optimisation task where the agent must safely land a spacecraft on a designated pad.
 - $\mathcal{X}$: $d_{\mathcal{X}} = 8$, including coordinates, linear velocities, angle, angular velocity, and boolean flags for leg ground contact.
 - $\mathcal{A}$: $d_{\mathcal{A}} = 2$, controls the main engine throttle and the side engine thrusters.
- **Hopper-v4**: A two-dimensional one-legged robot that aims to hop forward as fast as possible.
 - $\mathcal{X}$: $d_{\mathcal{X}} = 11$, comprising positional values and velocities of the body parts.
 - $\mathcal{A}$: $d_{\mathcal{A}} = 3$, controls the torques applied to the thigh, leg, and foot joints.
- **Walker2d-v5**: A two-dimensional bipedal robot that aims to walk forward.

- $\mathcal{X}$: $d_{\mathcal{X}} = 17$, including joint positions and velocities for both legs and the torso.
- $\mathcal{A}$: $d_{\mathcal{A}} = 6$, controls the torques for the joints.
- **HalfCheetah-v5**: A two-dimensional robot resembling a cheetah that aims to run forward.
 - $\mathcal{X}$: $d_{\mathcal{X}} = 17$, including positions and velocities of the torso and limbs.
 - $\mathcal{A}$: $d_{\mathcal{A}} = 6$, controls the torques on the forward and back thighs, shins, and feet.
- **Ant-v5**: A complex three-dimensional quadruped robot that aims to run forward.
 - $\mathcal{X}$: $d_{\mathcal{X}} = 27$, including the position, and orientation of the torso, joint angles, and velocities for all four legs.
 - $\mathcal{A}$: $d_{\mathcal{A}} = 8$, controls the torques on the hip and ankle joints of the four legs.

B.4.2. POLICIES

To construct the set of candidate policies Π , we trained agents using the Proximal Policy Optimization (PPO) algorithm (Schulman et al., 2017), as implemented in the Stable Baselines3 library (Raffin et al., 2021) (<https://github.com/DLR-RM/stable-baselines3>, MIT license). We utilised the standard `MlpPolicy` architecture. The agents were trained for a total of $T = 1 \times 10^6$ timesteps. The hyperparameters used for PPO training were consistent across all environments and are detailed in Table 3.

Table 3. PPO Hyperparameters used for training candidate policies.

Hyperparameter	Value
Total Timesteps (T)	1,000,000
Number of Environments (N_{envs})	8
Steps per Environment (N_{steps})	2048
Batch Size	1024
Learning Rate	3×10^{-4}
Discount Factor (γ)	0.97
GAE Lambda (λ)	0.95
Clip Range (ϵ)	0.2
Entropy Coefficient	0.01
Value Function Coefficient	0.5
Optimizer	Adam

We constructed the candidate policy set Π by saving checkpoints at evenly spaced intervals during training. Specifically, we selected policies at $\{0, 0.2T, 0.4T, 0.6T, 0.8T, T\}$ steps. The ground-truth preference ordering over Π , denoted by $\succ$, was established by estimating the expected discounted return $J(\pi)$ for each $\pi \in \Pi$. This was calculated via Monte Carlo estimation using 500 rollouts in the true environment.

B.4.3. FITTED Q EVALUATION

To obtain the proxy ground-truth rankings for decision-targeted training, we estimated the action-value function $Q^\pi(s, a)$ for each policy $\pi \in \Pi$ using Fitted Q Evaluation (FQE) (Le et al., 2019). Q^π_ψ was parametrised as an MLP with Layer Normalization and SiLU activations. The network takes the concatenated state and action vectors (s, a) as input and outputs a scalar Q-value.

To handle varying reward scales across environments and improve training stability, we standardised the rewards in the offline dataset to have zero mean and unit variance. Q^π_ψ was trained to predict the cumulative return of these standardised rewards. During inference, predictions were rescaled to the original magnitude.

Q^π_ψ were trained by minimising the expected Bellman error over the offline dataset $\mathcal{D}$. We employed a target network Q^π_ψ with an identical architecture, whose weights were updated via a hard update every 5 epochs.

We trained a separate Q^π_ψ for each policy in Π using the AdamW optimizer. To prevent gradient explosion, we applied gradient norm clipping. The specific hyperparameters are summarised in Table 4.

Table 4. Hyperparameters for Fitted Q Evaluation (FQE).

Hyperparameter	Value
Hidden Dimensions	[256, 256]
Activation Function	SiLU
Normalization	LayerNorm
Training Epochs	200
Batch Size	1,024
Learning Rate	3×10^{-4}
Target Update Frequency	Every 5 epochs
Target Action Samples (K)	32
Optimizer	AdamW
Gradient Clipping Norm	10.0

B.4.4. DIGITAL TWIN TRAINING

We evaluated our method using five distinct backbone architectures to demonstrate the architecture-agnostic nature of DT². All dynamics models were implemented as probabilistic ensembles (though we treat them as single stochastic models for the purpose of this work) that output the mean $\mu_\theta(\mathbf{x}, \mathbf{a})$ and diagonal log-variance $\log \Sigma_\theta(\mathbf{x}, \mathbf{a})$ of a Gaussian distribution $\mathcal{N}(\mu_\theta, \Sigma_\theta)$.

To improve training stability, we standardised inputs (state and action) to zero mean and unit variance using statistics computed from the offline dataset. The models were trained to predict the normalised state difference $\Delta \mathbf{x} = \mathbf{x}_{t+1} - \mathbf{x}_t$. The final prediction is obtained by denormalising the output and adding it to the current state.

To ensure numerical stability, the predicted log-variance is clamped to the range $[-10.0, 2.0]$. Additionally, during trajectory generation, the predicted next states are explicitly clipped to the specific physical bounds of the environment (e.g., joint limits) to ensure the simulation remains valid.

Base Architectures. We evaluated our method using five distinct feature extraction backbones. The output of each backbone is projected by two separate linear layers to produce the mean and log-variance vectors. The backbone details are as follows:

- **MLP:** A standard feed-forward network. It projects the input to the hidden dimension, followed by a SiLU (Swish) activation, a second linear layer, and another SiLU activation.
- **ResNet:** A residual network consisting of an input projection followed by 4 residual blocks. Each block applies: Linear $\rightarrow$ LayerNorm $\rightarrow$ SiLU $\rightarrow$ Dropout(0.1) $\rightarrow$ Linear $\rightarrow$ LayerNorm, with a skip connection adding the block input to the output. A final SiLU activation is applied after the blocks.
- **Neural ODE (NODE):** This backbone models the state evolution as an Ordinary Differential Equation (Chen et al., 2018). The hidden state transformation is defined as $\mathbf{z}(T) = \text{ODESolve}(\mathbf{z}(0), f, 0, T)$, where f is a neural network (Linear $\rightarrow$ Tanh $\rightarrow$ Linear) representing the derivative $d\mathbf{z}/dt$. We used an explicit Runge-Kutta 4 (RK4) solver with a fixed step size ($N = 1$ step).
- **Transformer:** Utilizes a standard TransformerEncoder (Vaswani et al., 2017) with 2 layers, 4 attention heads, and a feed-forward dimension of twice the hidden size. To capture temporal dependencies, the model receives a context window of the $T = 8$ most recent transitions. Learned positional embeddings are added to the input sequence.
- **GRU:** A Gated Recurrent Unit network (Cho et al., 2014) with 2 stacked layers. Similar to the Transformer, it processes a history sequence of length $T = 8$ to predict the next state.

All backbones used a hidden dimension of 256. The output of the backbone is projected by two separate linear layers to produce the mean and log-variance.

Table 5. Hyperparameters for Digital Twin Training (Standard and DT²).

Hyperparameter	Standard DT	DT ²
Optimizer	AdamW	AdamW
Learning Rate	3×10^{-4}	3×10^{-4}
Batch Size	1024	1024
Hidden Dimension	256	256
Gradient Clipping Norm	10.0	10.0
Max Epochs	2000	2000
Early Stopping Patience	20 epochs	20 epochs
<i>DT² Specific Parameters</i>		
Ranking Weight (λ)	N/A	0.1
Temperature (α)	N/A	1.0
Rollout Horizon (H)	N/A	20
Rollout Episodes (M)	N/A	128
Bootstrapping	N/A	Yes (using FQE)

Training Hyperparameters. All models were trained using the AdamW optimizer with gradient clipping. We employed early stopping based on the validation loss to prevent overfitting. The full set of hyperparameters is listed in Table 5.

N.B. in the Ant environment, we set $H = 50$. For DT² and standard DTs, we used $\mathcal{L}_{\text{sim}}^{\text{NLL}}$ as the simulation loss for the results displayed in §6.2, and results using $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ are in Appendix C.

VaGraM Baseline. We adapted the VaGraM method (Voelcker et al., 2022) to our offline ranking setting. Originally designed for active learning with a single policy, we extended it to weight transition errors based on the sensitivity of the value functions of all candidate policies in Π .

The model was trained to minimise a weighted Mean Squared Error (MSE) objective:

$$\mathcal{L}_{\text{VaGraM}}(\theta) = \mathbb{E}_{(\mathbf{x}, \mathbf{a}, \mathbf{x}') \sim \mathcal{D}} [||\mathbf{x}' - \mu_{\theta}(\mathbf{x}, \mathbf{a})||^2 \cdot w(\mathbf{x}')], \quad (29)$$

where the weight $w(\mathbf{x}')$ is the average squared norm of the value gradients at the next state across all policies:

$$w(\mathbf{x}') = \frac{1}{|\Pi|} \sum_{\pi \in \Pi} ||\nabla_{\mathbf{x}'} Q_{\psi}^{\pi}(\mathbf{x}', \pi(\mathbf{x}'))||_2^2. \quad (30)$$

The gradients $\nabla_{\mathbf{x}'} Q$ were computed via backpropagation through the pre-trained FQE networks (which were frozen during this process). We trained the VaGraM model using the MLP backbone and the same settings as the standard DTs.

B.5. Cancer Experiment

B.5.1. ENVIRONMENT DETAILS

We used the GhaffariCancerEnv from DTR-Bench/DTRGym (Luo et al., 2024) (<https://github.com/GilesLuo/DTR-Bench>, <https://github.com/GilesLuo/DTRGym>, MIT License), which simulates the treatment of cancer with metastasis under combined radiotherapy and chemotherapy. The environment is based on the mathematical model by Ghaffari et al. (2016).

- $\mathcal{X}$: $d_{\mathcal{X}} = 7$, log-transformed cell counts and drug concentrations.
- $\mathcal{A}$: $d_{\mathcal{A}} = 2$, radiotherapy dose ($D \in [0, 10]$ Gy), chemotherapy concentration ($v_M \in [0, 8]$).
- We utilized "Setting 5" of the environment, representing a challenging medical scenario. This includes:
 - State transition noise ($\sigma_{\text{state}} = 0.5$).
 - Observation noise ($\sigma_{\text{obs}} = 0.2$).
 - Pharmacokinetic/Pharmacodynamic parameter noise ($\sigma_{\text{pkpd}} = 0.1$).
 - A 50% probability of missing observations at any given timestep.

B.5.2. TRAINING POLICIES

For this experiment, we defined 5 interpretable clinical expert policies to construct Π_{train} :

1. **No Treatment:** Baseline policy with zero intervention ($D = 0, v_M = 0$).
2. **Fractionated Radiotherapy:** 2 Gy/day radiation, 5 days/week, with weekends off. No chemotherapy.
3. **Metronomic Chemotherapy:** Continuous low-dose chemotherapy ($v_M = 2.0$) administered daily. No radiotherapy.
4. **Adaptive Combined Therapy:** A reactive strategy that monitors tumour burden. If the log-size exceeds a threshold (50% of initial burden), it applies aggressive combination therapy ($D = 2.0, v_M = 4.0$). Otherwise, no treatment is given.
5. **Aggressive Combined Therapy:** High dose of both modalities ($D = 4.0, v_M = 6.0$) administered daily.

B.5.3. EVALUATION POLICIES

To test the generalisation capabilities of DT², we evaluated the models on a set of 11 unseen policies, including 6 PPO checkpoints along a training run of 200,000 steps, and the 5 following further interpretable policies:

- **Pulsed Chemotherapy:** High dose chemotherapy ($v_M = 5.0$) administered once every 21 days.
- **Hypofractionated Radiotherapy:** High radiation dose ($D = 8.0$) administered every 3 days.
- **Induction-Maintenance:** Aggressive combination therapy ($D = 2.0, v_M = 4.0$) for the first 30 days, followed by low-dose maintenance chemotherapy ($v_M = 1.5$).
- **Alternating Modality:** A 14-day cycle consisting of 7 days of daily radiotherapy ($D = 2.0$) followed by 7 days of daily chemotherapy ($v_M = 3.0$).
- **Dose Escalation:** Doses for both modalities linearly ramp up from 0 to a maximum ($D = 3.0, v_M = 5.0$) over the first 100 days.

B.5.4. DIGITAL TWIN TRAINING

The training pipeline followed the same structure as the continuous control experiments, with the following environment-specific adjustments:

- **Dataset:** We collected 1,000 steps from each of the 5 training expert policies.
- **Hyperparameters:** We used a batch size of 512 and a discount factor $\gamma = 0.99$. The Q-networks were trained for 500 epochs.

B.6. Lambda Ablation

For this experiment, we use the ResNet backbone in all the continuous control experiments. All settings are kept the same as those experiments, except we vary $\lambda \in \{0, 0.1, 0.25, 0.5, 0.75\}$.

C. Extended Results

We display results for the continuous control experiments using $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ for the simulation loss, rather than $\mathcal{L}_{\text{sim}}^{\text{NLL}}$ in Figure 6. We see similar takeaways from these results as those in §6.2. Namely, DT^2 outperforms its standard counterpart in 80% of the environment-architecture-metric couplings, and the best DT^2 model always outperforms VaGraM for regret and Spearman’s. In general, the DTs using $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ as their simulation loss seem to perform worse than when using $\mathcal{L}_{\text{sim}}^{\text{NLL}}$.




Figure 6. Regret and Spearman’s correlation for preference orderings from $\mathcal{L}_{\text{sim}}^{\text{MSE}}$ (hashed) and $\mathcal{L}_{\text{DT}^2}$ (solid) DTs across base architectures in six continuous control environments. Adapted VaGraM method shown in grey. We report averages over 10 seeds, with 95% CIs. For visual purposes, the top end of some error bars is cropped out.

D. Ranking Loss Ablation

We now consider some variants to the differentiable ranking loss used for DT². Let $\mathbf{y} \in \mathbb{R}^M$ be the vector of target proxy values (derived from FQE) for the M policies in a batch, and $\hat{\mathbf{y}} \in \mathbb{R}^M$ be the corresponding vector of cumulative returns predicted by the DT rollouts. Let $\mathcal{C} = \{(i, j) \mid 1 \leq i < j \leq M\}$ denote the set of all unique pairs in the batch.

Pairwise Hinge Loss. This variant enforces that the DT’s predicted value difference between two policies preserves the sign of the ground-truth difference, with a sufficient margin. To ensure the margin scales appropriately across different environments with varying reward magnitudes, we implemented an adaptive margin ϵ based on the mean absolute value of the targets:

$$\epsilon = 0.1 \cdot \frac{1}{M} \sum_{k=1}^M |y_k|. \quad (31)$$

The loss is defined as:

$$\mathcal{L}_{\text{Hinge}}(\theta) = \frac{1}{|\mathcal{C}|} \sum_{(i,j) \in \mathcal{C}} \max(0, \epsilon - \text{sign}(y_i - y_j)(\hat{y}_i - \hat{y}_j)). \quad (32)$$

This penalises the model if the predicted ordering is incorrect or if the separation between the pair is smaller than ϵ .

ListNet Loss. ListNet treats the ranking problem as a probability distribution matching problem. It maps the vector of values to a probability distribution using a softmax function.

Let P_y and $P_{\hat{y}}$ be the softmax distributions of the target and predicted values, respectively, controlled by a temperature parameter α :

$$P_y(\pi_i) = \frac{\exp(y_i/\alpha)}{\sum_{k=1}^M \exp(y_k/\alpha)}, \quad P_{\hat{y}}(\pi_i) = \frac{\exp(\hat{y}_i/\alpha)}{\sum_{k=1}^M \exp(\hat{y}_k/\alpha)}. \quad (33)$$

The loss is calculated as the Cross-Entropy between these distributions:

$$\mathcal{L}_{\text{ListNet}}(\theta) = - \sum_{i=1}^M P_y(\pi_i) \log(P_{\hat{y}}(\pi_i)). \quad (34)$$

Unlike the pairwise approaches, ListNet considers the entire list of policies simultaneously.

Continuous Control Results. We display results for the continuous control experiments with these variants in Figure 7, and report their average performances in Table 6. Across all environment-architecture-metric couplings, our smoothed Kendall formulation has the highest win-rate (43%), followed by ListNet (27%), Hinge (25%), and then the standard $\mathcal{L}_{\text{sim}}^{\text{NLL}}$ DT (5%). This indicates that including some form of ranking loss tends to improve performance, but our proposed method works best.

In Table 6, we can see that our smoothed Kendall formulation achieves the lowest average regret, and the highest average Spearman’s correlation. Interestingly, in this case, while the Hinge and ListNet formulations improve on $\mathcal{L}_{\text{sim}}^{\text{NLL}}$ in terms of average Spearman’s, they are worse in terms of average regret. As we can see from Figure 7, they each have some cases where they have a very poor regret (e.g. in the Pendulum and Lunarlander environments), which drives up their average regrets.

Table 6. Average Spearman’s correlation and regret for different DT variants across all architectures and continuous control task.

Loss	Spearman ($\uparrow$)	Regret ($\downarrow$)
$\mathcal{L}_{\text{sim}}^{\text{NLL}}$	0.47	6.21
Kendall	0.67	1.61
Hinge	0.62	8.22
ListNet	0.57	14.87




Figure 7. Regret and Spearman's correlation for preference orderings from $\mathcal{L}_{sim}^{NLL}$ DTs (hashed), DT^2 using our smoothed Kendall's loss (solid), DT^2 using a hinge loss (circles), and with DT^2 using a ListNet loss (dotted) across different base architectures. We report averages over 10 seeds, with 95% CIs. For visual purposes, the top end of some bars are cropped out.